


```
In [430]: plt.figure(figsize=(12, 6))
plt.title('Missing Values in Listings Dataset')
sns.barplot(df.isnull().sum().sort_values(ascending=False).reset_index(), x=
plt.xticks(rotation=90)
plt.xlabel('Features')
plt.ylabel('Number of Missing Values')

Out[430]: Text(0, 0.5, 'Number of Missing Values')
```



```
In [432] data_dict = pd.read_csv('listings_dictionary.csv', skiprows = 7).fillna('-')
In [433] data_dict.head()
```

Out[433]:

	Field	Type	Calculated	Description	Reference
0	id	integer	-	Airbnb's unique identifier for the listing	-
1	listing_url	text	y		-
2	scrape_id	bigint	y	Inside Airbnb "Scrape" this was part of	-
3	last_scraped	datetime	y	UTC. The date and time this listing was "scrap...	-
4	source	text	-	One of "neighbourhood search" or "previous scr...	-

In [434...

```
info_df = pd.concat([df.dtypes,df.isnull().sum()], axis = 1)
info_df = info_df.reset_index().rename({'index':'Field',0:'df dtype',1:'null
info_df['null_perc'] = info_df['null_perc']/len(df)
info_df = info_df.merge(data_dict, on = 'Field', how = 'inner')
```

In [435...

```
info_df.head()
```

Out[435]:

	Field	df dtype	null_perc	Type	Calculated	Description	Reference
0	id	int64	0.0	integer	-	Airbnb's unique identifier for the listing	-
1	listing_url	object	0.0	text	y		-
2	scrape_id	int64	0.0	bigint	y	Inside Airbnb "Scrape" this was part of	-
3	last_scraped	object	0.0	datetime	y	UTC. The date and time this listing was "scrap...	-
4	source	object	0.0	text	-	One of "neighbourhood search" or "previous scr...	-

Some columns need data type munging.

In [436...

```
info_df[info_df['null_perc']>0].sort_values('null_perc', ascending = False)
```

Out[436]:

	Field	df dtype	null_perc	Type	Calculated	Des
50	calendar_updated	float64	1.000000	date	-	
30	neighbourhood_group_cleansed	float64	1.000000	text	y	The neighbourhood as geocoded usir
8	neighborhood_overview	object	0.417124	text	-	Host's descriptio neighb
28	neighbourhood	object	0.417010	text	-	
15	host_about	object	0.345793	text	-	Description al
69	license	object	0.217535	text	-	licence/permit/reg
75	reviews_per_month	float64	0.214678	numeric	y	The average nu reviews per montl
68	review_scores_value	float64	0.214678	-	-	
60	first_review	object	0.214678	date	y	The da first/oldes
63	review_scores_accuracy	float64	0.214678	-	-	
64	review_scores_cleanliness	float64	0.214678	-	-	
65	review_scores_checkin	float64	0.214678	-	-	
66	review_scores_communication	float64	0.214678	-	-	
67	review_scores_location	float64	0.214678	-	-	
61	last_review	object	0.214678	date	y	The da last/newes
62	review_scores_rating	float64	0.214678	-	-	
14	host_location	object	0.194216	text	-	The host's self i
41	price	object	0.117741	currency	-	daily price currency;\nNO1
39	beds	float64	0.116941	integer	-	The number o
36	bathrooms	float64	0.113397	numeric	-	The number of ba in th
17	host_response_rate	object	0.081847	-	-	
16	host_response_time	object	0.081847	-	-	
18	host_acceptance_rate	object	0.051212	-	-	That rate at whic accepts bookin
22	host_neighbourhood	object	0.024577	text	-	
19	host_is_superhost	object	0.022634	boolean [t=true; f=false]	-	
38	bedrooms	float64	0.019776	integer	-	The number of be
7	description	object	0.018633	text	-	Detailed descr th
51	has_availability	object	0.011317	boolean	-	[t=true; the
49	maximum_nights_avg_ntm	float64	0.001943	numeric	y	maximum_nig

						from the
48	minimum_nights_avg_ntm	float64	0.001943	numeric	y	the minimum_nig from the
47	maximum_maximum_nights	float64	0.001943	integer	y	the maximum_nig from the
46	minimum_maximum_nights	float64	0.001943	integer	y	the maximum_nig from th
44	minimum_minimum_nights	float64	0.001943	integer	y	the minimum_nig from th
45	maximum_minimum_nights	float64	0.001943	integer	y	the minimum_nig from the
37	bathrooms_text	object	0.001486	string	-	The number of ba in the listing.

Calendar updated and neighborhood group cleaned are both columns with all null values. Other columns with over 20% missing values are concerning and should be investigated further.

Data Preprocessing

Dropping Columns

Columns that do not contribute any information to our project scope are dropped prior to any further data cleaning:

- scrape_id, last_scraped, source, listing_url, calendar_updated, calendar_last_scraped
 - stationary values as the data does not go back further than one scrape
- name, picture_url, host_url, host_thumbnail_url, host_picture_url, host_id, host_location, host_since, host_about, host_response_time, host_response_rate, host_acceptance_rate, host_is_superhost, host_total_listings_count, host_verifications, host_name, host_has_profile_pic, host_identity_verified, calculated_host_listings_count, calculated_host_listings_count_entire_homes, calculated_host_listings_count_private_rooms, calculated_host_listings_count_shared_rooms, first_review, last_review
 - irrelevant to the project
- host_neighborhood, neighborhood, neighbourhood_group_cleaned, host_listings_count, reviews_per_month, property_type, neighbourhood_cleaned, neighborhood_overview
 - other related or calculated column that provides pertinent information is included
 - Latitude and Longitude will be used to provide distance feature
- min_min, min_max, max_min, maxmax, availability, has_availability, availability_eoy, estimated_occupancy_1365d, estimated_revenue_1365d
 - looks into the future, and therefore, dropped to minimize data leakage

```
In [437]: drop_col = ['listing_url', 'scrape_id', 'last_scraped', 'source', 'calendar_l',
                  'name', 'picture_url', 'host_url', 'host_thumbnail_url', 'host_pi',
                  'host_response_time', 'host_response_rate', 'host_acceptance_ra',
                  'host_total_listings_count', 'host_verifications',
                  'host_has_profile_pic', 'host_identity_verified',
                  'host_neighbourhood', 'host_listings_count', 'neighbourhood', 'h',
                  'neighbourhood_group_cleaned', 'neighbourhood_cleaned', 'revi',
                  'minimum_nights', 'maximum_nights', 'minimum_minimum_nights',
                  'maximum_minimum_nights', 'minimum_maximum_nights',
                  'maximum_maximum_nights', 'minimum_nights_avg_ntm',
                  'maximum_nights_avg_ntm', 'availability_30', 'availability_60',
                  'availability_365', 'has_availability', 'availability_eoy', 'esti',
                  'calculated_host_listings_count',
                  'calculated_host_listings_count_entire_homes',
                  'calculated_host_listings_count_private_rooms',
                  'calculated_host_listings_count_shared_rooms']
```

```
In [438... df = df.drop(drop_col, axis = 1)
info_df = info_df[~info_df['Field'].isin(drop_col)]
```

```
In [439... info_df[info_df['null_perc']>0].sort_values('null_perc', ascending = False)
```

Out[439]:

	Field	df dtype	null_perc	Type	Calculated	Description
69	license	object	0.217535	text	-	licence/permit/registration number
62	review_scores_rating	float64	0.214678	-	-	
63	review_scores_accuracy	float64	0.214678	-	-	
64	review_scores_cleanliness	float64	0.214678	-	-	
65	review_scores_checkin	float64	0.214678	-	-	
66	review_scores_communication	float64	0.214678	-	-	
67	review_scores_location	float64	0.214678	-	-	
68	review_scores_value	float64	0.214678	-	-	
41	price	object	0.117741	currency	-	daily price in currency. NOTE that some values are in US dollars and some are in the local currency.
39	beds	float64	0.116941	integer	-	The number of beds in the property.
36	bathrooms	float64	0.113397	numeric	-	The number of bathrooms in the property.
38	bedrooms	float64	0.019776	integer	-	The number of bedrooms in the property.
7	description	object	0.018633	text	-	Detailed description of the property.
37	bathrooms_text	object	0.001486	string	-	The number of bathrooms in the listing. If the listing has multiple bathrooms, this is the total number.

Dropping Rows

It appears that some properties are new and are missing review information (about 21.5%). These properties will be removed prior to further analysis.

Additionally, properties that do not have a description will also be removed. It is a key feature we will be using NLP on, and there are less than 2% of properties without a description.

```
In [440... df = df[(df['review_scores_rating'].notnull()) & (df['description'].notnull())]
```

Data Type Munging

```
In [441... df['price'] = df['price'].str.replace(r'[$,]', '', regex = True).astype(float)
```

Cleaning bathrooms_text

```
In [442... df['bathrooms_text'].str.split().str[-1].value_counts()
```

Out[442]:

bath	4357
baths	2449
Half-bath	3
half-bath	2
Name: bathrooms_text, dtype: int64	

```
In [443... df[df['bathrooms_text'].str.contains('share', na = False, case = False)]['bathrooms_text']
```

Out[443]:

bath	589
baths	325
half-bath	1
Name: bathrooms_text, dtype: int64	

```
In [444... df[df['bathrooms_text'].str.contains('half-bath', na = False, case = False)]
```

Out[444]:

	bathrooms	bathrooms_text
401	0.5	Shared half-bath
1927	NaN	Half-bath
5720	0.5	Half-bath
5914	0.5	Half-bath
6479	0.5	Private half-bath

```
In [445... df['bathrooms_text'].str.split().str[0].unique()
```

Out[445]:

array(['1', '2', '1.5', '3', '2.5', '3.5', '11', nan, '4', '0', 'Shared', '4.5', '5', '11.5', '6.5', '5.5', '7', 'Half-bath', '6', '9.5', '7.5', '8', '8.5', '9', 'Private'], dtype=object)

```
In [446... df['bathrooms_text'].str.split().str[1].unique()
```

Out[446]:

array(['shared', 'bath', 'baths', 'private', nan, 'half-bath'], dtype=object)

It appears there are shared baths. Check if there are other categories of bath rooms.
Assume if not listed as shared, it is private.

```
In [447... df['bathrooms_shared'] = np.where(df['bathrooms_text'].str.contains('share',  
df = df.drop('bathrooms_text', axis = 1)
```

Replacing license to a categorical variable

```
In [448]: pd.Series(np.where(df['license'].fillna('null').str.replace(r'[A-Za-z]', '', 'Licensed',df['license'])).value_counts())
```

```
Out[448]: Licensed
5872
City registration pending
178
32+ Days Listing
61
32+days Listing
27
City Registration Pending
7
Registered
4
Per city of chicago, no registration # is needed since this rental is 32 da
ys or more.      1
DOB-111617
1
City registration pending R19000048093
1
Applied for registration
1
Chicago registration number pending
1
2120298, 2120297
1
Registration number pending
1
Registration pending
1
c
1
Pending
1
PENDING
1
City registration permit pending
1
city registration pending
1
47-5611763
1
pending
1
dtype: int64
```

```
In [449]: pd.Series(np.where((df['license'].str.contains('32', na = False))&\
(df['license'].str.contains('day', na = False, case = False)), 'NA',
np.where((df['license'].str.contains('pending', na = False, case =
(df['license'].str.contains('applied', na = False, case = False)),
np.where((df['license'].str.contains(r'\d', regex = True))\
(df['license'].str.contains('registered', na = Fa
```

```
Out[449]: LICENSE      5879
nan              660
PEND             195
NA               89
dtype: int64
```

```
In [450]: df['license'] = np.where((df['license'].str.contains('32', na = False))&\
(df['license'].str.contains('day', na = False, case = False)), 'NA',
np.where((df['license'].str.contains('pending', na = False, case =
(df['license'].str.contains('applied', na = False, case = False)),
np.where((df['license'].str.contains(r'\d', regex = True))\
(df['license'].str.contains('registered', na = Fa
```

Creating Geo-Features

- Distance from the Bean

```
In [451]: def calculate_to_loc(dataframe, fixed_point):
return dataframe.apply(
lambda row: geodesic((row['latitude'], row['longitude']), fixed_poi
axis=1
)
```

```
In [452]: df['latitude'].isnull().sum(), df['longitude'].isnull().sum()
```

```
Out[452]: (0, 0)
```

```
In [453]: fixed_point = (41.892423, -87.634049) # Bean (Downtown Chicago)
fixed_lat = 41.892423
fixed_lon = -87.634049

df['km_DT'] = calculate_to_loc(df, fixed_point)
df['lat_diff'] = df['latitude'] - fixed_lat
df['lon_diff'] = df['longitude'] - fixed_lon
```

```
In [454]: df.drop(['latitude', 'longitude'], axis = 1, inplace = True)
```

Free Text Fields

- Description, amenities

```
In [455]: df.head().T
```

Out[455]:

		0	1	2	3	4
	id	2384	7126	10945	28749	71930
	description	Solo Hyde Park visitors are invited stay in th...	A very small studio in a wonderful neighborhood.	Beautiful first floor apartment in Historic Ol...	Located on a peaceful treelined street in ener...	A peaceful sharec space in Chicago's Ukrainian...
	room_type	Private room	Entire home/apt	Entire home/apt	Entire home/apt	Private room
	accommodates	1	2	4	6	2
	bathrooms	1.0	1.0	1.0	2.0	1.0
	bedrooms	1.0	1.0	2.0	3.0	1.0
	beds	1.0	1.0	2.0	3.0	1.0
	amenities	["Host greets you", "Hot water kettle", "Carbo...	["Window AC unit", "Central heating", "Dishes ...	["Window AC unit", "Dishes and silverware", "H...	["Dishes and silverware", "TV with DVD player,...	["Dishes and silverware", "Dedicated workspace...
	price	125.0	81.0	187.0	196.0	76.0
	number_of_reviews	250	569	117	244	125
	number_of_reviews_ltm	16	53	34	47	15
	number_of_reviews_l30d	0	0	0	3	0
	number_of_reviews_ly	20	52	36	42	20
	review_scores_rating	4.99	4.72	4.72	4.82	4.85
	review_scores_accuracy	4.98	4.85	4.83	4.87	4.90
	review_scores_cleanliness	4.99	4.57	4.81	4.76	4.75
	review_scores_checkin	4.99	4.91	4.83	4.94	4.96
	review_scores_communication	4.98	4.88	4.87	4.88	4.95
	review_scores_location	4.95	4.9	4.97	4.93	4.84
	review_scores_value	4.94	4.76	4.72	4.72	4.84
	license	LICENSE	LICENSE	LICENSE	LICENSE	LICENSE
	instant_bookable	f	f	t	f	
	bathrooms_shared	1	0	0	0	
	km_DT	12.2284	3.965647	2.222031	6.297014	3.781326
	lat_diff	-0.104523	0.009237	0.019537	0.027127	0.003727
	lon_diff	0.046249	-0.046161	-0.005761	-0.066641	-0.04529

Removing Emojis

```
In [456]: def contains_emoji(text):
    if pd.isna(text):
        return 0
    text = str(text)
    emoji_pattern = re.compile("[
        u"\U0001F600-\U0001F64F" # emoticons
        u"\U0001F300-\U0001F5FF" # symbols & pictogr
        u"\U0001F680-\U0001F6FF" # transport & map s
        u"\U0001F1E0-\U0001F1FF" # flags (iOS)
        u"\U00002700-\U000027BF" # Dingbats
        u"\U0001F900-\U0001F9FF" # Supplemental Sym
        u"\U0001F700-\U0001F77F" # Alchemical Symbol
        u"\U0001F780-\U0001F7FF" # Geometric Shapes
        u"\U0001F800-\U0001F8FF" # Supplemental Arroc
        u"\U0001F100-\U0001F1FF" # Enclosed Alphanum
        u"\U0001F200-\U0001F2FF" # Enclosed Ideograp
        u"\U0001F300-\U0001F5FF" # Miscellaneous Sym
        u"\U0001F600-\U0001F64F" # Emoticons
        u"\U0001F650-\U0001F67F" # Ornamental Dingba
        u"\U0001F680-\U0001F6FF" # Transport and Map
        u"\U0001F700-\U0001F77F" # Alchemical Symbol
        u"\U0001F780-\U0001F7FF" # Geometric Shapes
        u"\U0001F800-\U0001F8FF" # Supplemental Arroc
        u"\U0001F900-\U0001F9FF" # Supplemental Sym
        u"\U0001FA00-\U0001FA6F" # Chess Symbols
        u"\U0001FA70-\U0001FAFF" # Symbols and Pict
    ]+", flags=re.UNICODE)
    return 1 if emoji_pattern.search(text) else 0

def remove_emojis(text):
    if pd.isna(text):
        return text
    text = str(text)
    emoji_pattern = re.compile("[
        u"\U0001F600-\U0001F64F" # emoticons
        u"\U0001F300-\U0001F5FF" # symbols & pictogr
        u"\U0001F680-\U0001F6FF" # transport & map s
        u"\U0001F1E0-\U0001F1FF" # flags (iOS)
        u"\U00002700-\U000027BF" # Dingbats
        u"\U0001F900-\U0001F9FF" # Supplemental Sym
        u"\U0001F700-\U0001F77F" # Alchemical Symbol
        u"\U0001F780-\U0001F7FF" # Geometric Shapes
        u"\U0001F800-\U0001F8FF" # Supplemental Arroc
        u"\U0001F900-\U0001F9FF" # Supplemental Sym
        u"\U0001FA00-\U0001FA6F" # Chess Symbols
    ]+", flags=re.UNICODE)
    return text.replace(emoji_pattern, "")
```

```

        u"\U0001FA70-\U0001FAFF" # Symbols and Pictographs
    ]+", flags=re.UNICODE)
    return emoji_pattern.sub(r'', text)

```

```

In [457]: df['emoji_description'] = df['description'].apply(contains_emoji).values

```

```

In [458]: df['description'] = df['description'].apply(remove_emojis)

```

Feature Engineering (NLP)

Key Phrase and Keyword Extraction

- Using brute force (occurrence) + RAKE algorithm to identify top keywords and key phrases.
- Uses levenshtein distance to keep "unique" terms.
- Vectorize description field.

```

In [459]: def preprocess_text(text):
    text = text.lower()
    # remove HTML tags
    text = re.sub(r'<[>]+>', ' ', text)
    # remove special characters
    text = re.sub(r'[\W\s]', ' ', text)
    # remove extra spaces
    text = re.sub(r'\s+', ' ', text).strip()
    return text

stop = set(nltk.corpus.stopwords.words('english'))

# Extract keywords from text: using brute count
def extract_keywords_count(text):
    words = text.split()
    # filter keywords for at least > 4 characters and not a stop word
    words = list(filter(lambda x: len(x) >= 4 and x not in stop, words))
    # only top 10 occurrences
    keywords = [element[0] for element in Counter(words).most_common(10)]
    return keywords

# Getting the key phrases using RAKE Algorithm
def extract_key_phrase(text):
    key_phrases = Rake()
    key_phrases.extract_keywords_from_text(preprocess_text(text))
    ranked_phrases = key_phrases.get_ranked_phrases()
    ranked_phrases = list(dict.fromkeys(ranked_phrases))[:10]
    return ranked_phrases

# Unique between key words brute count + keyphrases RAKE
def merge_description(lst):
    if len(lst) < 3:
        return lst
    filtered_list = [lst[0]]
    for item in lst[1:]:
        distance_score = 1 - (levenshtein_distance(item.replace(' ', ''), lst[0]))
        if distance_score <= .5:
            filtered_list.append(item)
    return filtered_list

```

```

In [460]: df['description_1'] = df['description'].map(preprocess_text).map(extract_key)
df['description_2'] = df['description'].map(preprocess_text).map(extract_key)
df['description_v'] = df.apply(lambda row: merge_description(list(set(row['description_1'], 'description_2'])), axis=1, inplace=True)

```

NMF to cluster the description keywords

```

In [461]: tfidf_vectorizer = TfidfVectorizer(max_df=0.85, min_df=2, stop_words='english')
tfidf_matrix = tfidf_vectorizer.fit_transform(df['description_v'])


```

Grid Search Number of Clusters


```
In [406... reconstruction_errors = []
components_range = range(2, 100, 5)

for n_comp in components_range:
    nmf_model_eval = NMF(n_components=n_comp, random_state=1, init='nndsvda')
    nmf_model_eval.fit(tfidf_matrix)
    reconstruction_errors.append(nmf_model_eval.reconstruction_err_)
    print(f" N_components: {n_comp}, Reconstruction Error: {nmf_model_eval.
plt.figure(figsize=(10, 6))
plt.plot(components_range, reconstruction_errors, marker='o', linestyle='-')
plt.title('NMF Reconstruction Error vs. Number of Components (Elbow Method)')
plt.xlabel('Number of Components (n_components)')
plt.ylabel('Reconstruction Error (Frobenius Norm)')
plt.xticks(list(components_range)) # Ensure all component numbers are shown
plt.grid(True)
plt.show()
```

```
N_components: 2, Reconstruction Error: 80.1541
N_components: 7, Reconstruction Error: 78.1745
N_components: 12, Reconstruction Error: 77.0799
N_components: 17, Reconstruction Error: 76.1293
N_components: 22, Reconstruction Error: 75.2709
N_components: 27, Reconstruction Error: 74.4900
N_components: 32, Reconstruction Error: 73.7519
N_components: 37, Reconstruction Error: 73.0771
N_components: 42, Reconstruction Error: 72.3759
N_components: 47, Reconstruction Error: 71.7331
N_components: 52, Reconstruction Error: 71.1636
N_components: 57, Reconstruction Error: 70.5129
N_components: 62, Reconstruction Error: 69.9972
N_components: 67, Reconstruction Error: 69.4567
N_components: 72, Reconstruction Error: 68.9205
N_components: 77, Reconstruction Error: 68.3953
N_components: 82, Reconstruction Error: 67.8392
N_components: 87, Reconstruction Error: 67.3672
N_components: 92, Reconstruction Error: 66.9096
N_components: 97, Reconstruction Error: 66.4651
```



While there is no clear "elbow", there's a significant improvement going from 2 to 12 components. The decrease from 12 to 17 is still noticeable (0.9506). After 17 components, the decreases become more consistent and smaller (generally below 0.8 and moving towards 0.4-0.5).

```
In [462... n_cluster = 17
nmf = NMF(n_components=n_cluster, random_state=1, init='nndsvda', max_iter=5)
nmf_w = nmf.fit_transform(tfidf_matrix) # Document-topic matrix
nmf_H = nmf.components_ # Topic-word matrix

def display_topics(model, feature_names, no_top_words):
    for topic_idx, topic in enumerate(model.components_):
        print(f"Topic {topic_idx + 1}:")
        print(" ".join([feature_names[i]
                        for i in topic.argsort()[::-no_top_words - 1:-1]]))
    print("\n")

# Get the feature names (words) from the TF-IDF vectorizer
tfidf_feature_names = tfidf_vectorizer.get_feature_names_out()

# Display the top words for each topic
num_top_words = 5 # Number of top words to display per topic
print(f"Top {num_top_words} words per topic:\n")
display_topics(nmf, tfidf_feature_names, num_top_words)
```

Top 5 words per topic:

```
Topic 1:
minutes away downtown chicago drive
Topic 2:
blueground home start living love
Topic 3:
centrally place located peaceful simple
Topic 4:
special rates message queen building
Topic 5:
room private shared bed living
Topic 6:
stylish space perfect experience vibrant
Topic 7:
walking chicago distance field attractions
Topic 8:
great place time comfortable family
Topic 9:
square logan blue line garden
Topic 10:
studio west away shops blue
Topic 11:
night river north level separate
Topic 12:
unit bedroom condo bath newly
Topic 13:
park lincoln wicker hyde grant
Topic 14:
access easy enjoy group perfectly
Topic 15:
walk minute min line wrigley
Topic 16:
free parking street quiet safe
Topic 17:
apartment bedroom fully furnished living
```

```
In [463]: topic_feature_names = [f'nmf_topic_{i}' for i in range(n_cluster)]

# Create a DataFrame from nmf_W with the new column names
nmf_df = pd.DataFrame(nmf_W, columns=topic_feature_names, index=df.index)

# Concatenate the original DataFrame with the new NMF features DataFrame
df = pd.concat([df, nmf_df], axis=1)
```

```
In [464]: df.drop(['description_v', 'description'], axis = 1, inplace = True)
```

K-Means Clustering Amenities

```
In [465]: df['amenities'] = df['amenities'].apply(ast.literal_eval).apply(lambda x: '
tfidf_vectorizer = TfidfVectorizer(max_df=0.85, min_df=2, stop_words='englis
tfidf_matrix = tfidf_vectorizer.fit_transform(df['amenities'])
```

```
In [411]: inertia_values = []
silhouette_scores = []
components_range = range(2, 21)

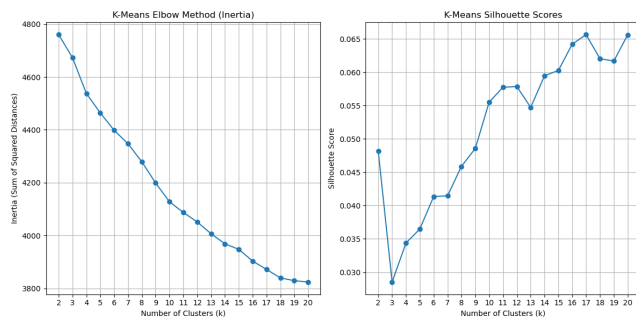
for n_comp in components_range:
    kmeans_model = KMeans(n_clusters=n_comp, random_state=1, n_init='auto')
    # Fit the model to the TF-IDF matrix
    kmeans_model.fit(tfidf_matrix)
    # Store the inertia value
    inertia_values.append(kmeans_model.inertia_)
    # Predict clusters for the current K
    cluster_labels = kmeans_model.predict(tfidf_matrix)
    # Calculate silhouette score
    score = silhouette_score(tfidf_matrix, cluster_labels)
    silhouette_scores.append(score)
    print(f" K: {n_comp}, Inertia: {kmeans_model.inertia_:.2f}, Silhouette
```

```
K: 2, Inertia: 4760.26, Silhouette Score: 0.048
K: 3, Inertia: 4673.09, Silhouette Score: 0.029
K: 4, Inertia: 4537.48, Silhouette Score: 0.034
K: 5, Inertia: 4464.52, Silhouette Score: 0.036
K: 6, Inertia: 4398.31, Silhouette Score: 0.041
K: 7, Inertia: 4347.70, Silhouette Score: 0.041
K: 8, Inertia: 4279.05, Silhouette Score: 0.046
K: 9, Inertia: 4198.69, Silhouette Score: 0.049
K: 10, Inertia: 4127.57, Silhouette Score: 0.055
K: 11, Inertia: 4087.27, Silhouette Score: 0.058
K: 12, Inertia: 4051.29, Silhouette Score: 0.058
K: 13, Inertia: 4006.76, Silhouette Score: 0.055
K: 14, Inertia: 3968.52, Silhouette Score: 0.059
K: 15, Inertia: 3947.68, Silhouette Score: 0.060
K: 16, Inertia: 3903.22, Silhouette Score: 0.064
K: 17, Inertia: 3871.57, Silhouette Score: 0.066
K: 18, Inertia: 3839.72, Silhouette Score: 0.062
K: 19, Inertia: 3829.00, Silhouette Score: 0.062
K: 20, Inertia: 3824.01, Silhouette Score: 0.066
```

```
In [412]: plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(components_range, inertia_values, marker='o', linestyle='-')
plt.title('K-Means Elbow Method (Inertia)')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia (Sum of Squared Distances)')
plt.xticks(list(components_range))
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(components_range, silhouette_scores, marker='o', linestyle='-')
plt.title('K-Means Silhouette Scores')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Silhouette Score')
plt.xticks(list(components_range))
plt.grid(True)

plt.tight_layout()
plt.show()
```



Elbow Method:

- K=2 to K=4: The inertia drops significantly (e.g., from 4760.26 to 4537.48, a drop of over 220 in two steps).
- K=4 to K=10: The drops continue but become somewhat less steep (e.g., from ~70 to ~80 per step of 1).
- K=10 to K=15: The reduction in inertia becomes noticeably smaller, hovering around 20-40 per step.
- K=15 to K=20: The drops become quite small, especially after K=18 (a drop of only 10.72) and K=19 (a drop of only 4.99).

Silhouette: The scores start low and generally increase, but the overall values remain quite low (all below 0.1). This is common for text data, but it suggests the clusters might not be extremely distinct or well-separated. The score reaches a local high at K=10 (0.055). It then continues to gradually climb, reaching its highest observed values at K=17 (0.066) and K=20 (0.066).

K=17 is a strong candidate because it falls within the region where the inertia curve starts to flatten, and it achieves one of the highest silhouette scores.

```
In [466] kmeans_model = KMeans(n_clusters=17, random_state=1, n_init='auto')
cluster_assignments = kmeans_model.fit_predict(tfidf_matrix)
df['amenities_KMeans'] = cluster_assignments

In [467] df.drop(['amenities'], axis = 1, inplace = True)
```

Feature Engineering (Encoding Categorical Variables)

```
In [472] df = pd.get_dummies(df, columns = ['room_type', 'license', 'instant_bookable'],
```

Scaling Continuous Data

- Min Max Scaler was used to keep features at 0-1 unit

```
In [473] cont_cols = ['km_pt', 'lat_diff', 'lon_diff', 'bathrooms', 'bedrooms', 'beds', 'ac',
'number_of_reviews_130d', 'number_of_reviews_ly', 'review_scores_rating',
'review_scores_accuracy', 'review_scores_cleanliness',
'review_scores_checkin', 'review_scores_communication',
'review_scores_location', 'review_scores_value']

In [ ]: scaler = MinMaxScaler()
continuous = scaler.fit_transform(df[cont_cols])
df[cont_cols] = continuous

In [480] df = df.set_index('id')

In [ ]: df.to_pickle('data.pkl')
df = pd.read_pickle('data.pkl')
```

Model Building

- Agglomerative Clustering
 - Does not require the number of clusters to be specified in advance
 - Bottoms up processing
 - Small data set does not get computationally expensive
 - Properties exhibit hierarchical relationships

```
In [490] len(df.dropna()), len(df)
Out[490]: (6077, 6823)
```

Since there are not many missing NaNs, we will drop rows that contain a null value.

```
In [491] df = df.dropna()
```

Dendrogram

In [503]: `X = df.values`

Euclidean distance is the most common and intuitive distance metric. It measures the "straight-line" distance between two points in a multi-dimensional space.

Since the features are all numerical and have been Min-Max scaled to a 0-1 range, Euclidean distance becomes very meaningful.

Scaling ensures that no single feature dominates the distance calculation merely because it has a larger range of values. It implies that proximity in this feature space is a direct measure of similarity.

Ward linkage tends to produce clusters that are roughly spherical, compact, and of similar size. It's generally robust and less susceptible to noise compared to other linkage methods like 'single' linkage. It tries to keep the clusters as homogeneous as possible by minimizing the increase in total within-cluster variance.

```
In [ ]: Z = linkage(X, method='ward', metric='euclidean')

plt.figure(figsize=(15, 7))
plt.title('Hierarchical Clustering Dendrogram for All Features')
plt.xlabel('Sample Index or (Cluster Size)')
plt.ylabel('Distance (Ward Linkage)')
dendrogram(
    Z,
    truncate_mode='lastp', # show only the last p merged clusters
    p=30, # show only the last 30 merges
    leaf_rotation=90,
    leaf_font_size=8,
    show_leaf_counts=True,
    above_threshold_color='blue',
)
```

```
Out [ ]: {'icoord': [[25.0, 25.0, 35.0, 35.0],
[15.0, 15.0, 30.0, 30.0],
[45.0, 45.0, 55.0, 55.0],
[22.5, 22.5, 50.0, 50.0],
[5.0, 5.0, 36.25, 36.25],
[65.0, 65.0, 75.0, 75.0],
[85.0, 85.0, 95.0, 95.0],
[105.0, 105.0, 115.0, 115.0],
[125.0, 125.0, 135.0, 135.0],
[155.0, 155.0, 165.0, 165.0],
[145.0, 145.0, 160.0, 160.0],
[175.0, 175.0, 185.0, 185.0],
[215.0, 215.0, 225.0, 225.0],
[205.0, 205.0, 220.0, 220.0],
[195.0, 195.0, 212.5, 212.5],
[180.0, 180.0, 203.75, 203.75],
[235.0, 235.0, 245.0, 245.0],
[285.0, 285.0, 295.0, 295.0],
[275.0, 275.0, 290.0, 290.0],

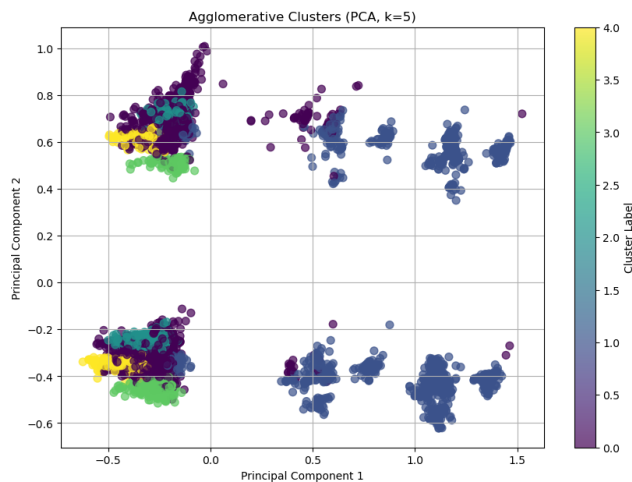
[265.0, 265.0, 282.5, 282.5],
[255.0, 255.0, 273.75, 273.75],
[240.0, 240.0, 264.375, 264.375],
[191.875, 191.875, 252.1875, 252.1875],
[152.5, 152.5, 222.03125, 222.03125],
[130.0, 130.0, 187.265625, 187.265625],
[110.0, 110.0, 158.6328125, 158.6328125],
[90.0, 90.0, 134.31640625, 134.31640625],
[70.0, 70.0, 112.158203125, 112.158203125],
[20.625, 20.625, 91.0791015625, 91.0791015625]],
'dcoord': [[0.0, 15.151337945036339, 15.151337945036339, 0.0],
[0.0, 16.408342949409203, 16.408342949409203, 15.151337945036339],
[0.0, 18.36196818367173, 18.36196818367173, 0.0],
[16.408342949409203,
21.570525473978876,
21.570525473978876,
18.36196818367173],
[0.0, 23.950096881173756, 23.950096881173756, 21.570525473978876],
[0.0, 15.992429147952588, 15.992429147952588, 0.0],
[0.0, 13.212212382287257, 13.212212382287257, 0.0],
[0.0, 13.655983565268999, 13.655983565268999, 0.0],
[0.0, 14.44988033264926, 14.44988033264926, 0.0],
[0.0, 13.515377619159258, 13.515377619159258, 0.0],
[0.0, 21.22031430411815, 21.22031430411815, 13.515377619159258],
[0.0, 17.43670841951672, 17.43670841951672, 0.0],
[0.0, 13.995079906684452, 13.995079906684452, 0.0],
[0.0, 17.036356426258845, 17.036356426258845, 13.995079906684452],
[0.0, 19.736068534774237, 19.736068534774237, 17.036356426258845],
[17.43670841951672,
20.97872637659182,
20.97872637659182,
19.736068534774237],
[0.0, 12.877101612671687, 12.877101612671687, 0.0],
[0.0, 13.738732587920051, 13.738732587920051, 0.0],
[0.0, 15.059311979356496, 15.059311979356496, 13.738732587920051],
[0.0, 16.171825698172135, 16.171825698172135, 15.059311979356496],
[0.0, 18.923958171748076, 18.923958171748076, 16.171825698172135],
[12.877101612671687,
24.023390578785133,
24.023390578785133,
18.923958171748076],
[20.97872637659182,
25.008186373223584,
25.008186373223584,
24.023390578785133],
[21.22031430411815,
27.35395968429988,
27.35395968429988,
25.008186373223584],
[14.44988033264926,
30.132718068034876,
30.132718068034876,
27.35395968429988],
[13.655983565268999,
31.142268355200095,
31.142268355200095,
30.132718068034876],
```


Hierarchical Clustering Dendrogram for All Features

The dendrogram illustrates the hierarchical clustering of 32 samples (labeled c1462 to c1333) based on distance. The y-axis represents the Distance (Ward Linkage) from 0 to 50. The x-axis represents the Sample Index (or Cluster Size). The dendrogram shows the hierarchical merging of clusters, with a blue line highlighting the final merge of clusters c1462-c1484 and c1385-c1407 at a distance of approximately 51. Other clusters are highlighted in orange and green.

```
In [510]: pca = PCA(n_components=2, random_state=1)
pca_components = pca.fit_transform(X)

plt.figure(figsize=(10, 7))
plt.scatter(pca_components[:, 0], pca_components[:, 1], c=labels, cmap='viridis')
plt.title('Agglomerative Clusters (PCA, k={n_clusters})')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster Label')
plt.grid(True)
```



Cluster Separation and Distribution:

Vertical Separation by PC2: There's a very clear and strong separation along Principal Component 2 (the y-axis). The clusters tend to group into two distinct horizontal bands of -0.2 to -0.6 and 0.4 to 1.0.

The yellow and green clusters are quite compact and visually well-separated from the large dark purple and blue clusters. The teal clusters are also moderately well-separated.

- Within the top band, we see distinct yellow, green, and a teal cluster.
- Similarly, within the bottom band, we see another set of yellow, green, and a teal cluster.

GridSearch N Clusters

```
In [512]: n_clusters_to_test = [3, 5, 6, 8, 9, 10]

results = []

for n_clusters in n_clusters_to_test:
    model = AgglomerativeClustering(n_clusters=n_clusters, metric='euclidean')
    labels = model.fit_predict(X)
    silhouette_avg = silhouette_score(X, labels)
    davies_bouldin_idx = davies_bouldin_score(X, labels)
    calinski_harabasz_idx = calinski_harabasz_score(X, labels)
    # Store results
    results.append({
        'n_clusters': n_clusters,
        'silhouette_score': silhouette_avg,
        'davies_bouldin_index': davies_bouldin_idx,
        'calinski_harabasz_index': calinski_harabasz_idx
    })

results_df = pd.DataFrame(results)
```

```
In [513]: results_df
```

```
Out[513]:
```

	n_clusters	silhouette_score	davies_bouldin_index	calinski_harabasz_index
0	3	0.078284	2.029176	648.675525
1	5	0.138328	1.888306	558.849108
2	6	0.168462	1.878533	550.503589
3	8	0.196052	2.048123	524.217370
4	9	0.206415	1.951481	514.261610
5	10	0.213882	1.965713	510.894984

Trade-off at K=6:

This gives us the lowest Davies-Bouldin Index (1.878), indicating the best balance of compactness and separation. The Silhouette Score is also reasonably good for this dataset (0.168), although not the absolute highest. The Calinski-Harabasz index is still relatively high compared to higher K values, suggesting it retains some overall cluster density.

```
In [514]: n_clusters = 6
base_model = AgglomerativeClustering(n_clusters=n_clusters, metric='euclidean')
base_model.fit(X)
labels = base_model.labels_


silhouette_avg = silhouette_score(X, labels)
davies_bouldin_idx = davies_bouldin_score(X, labels)
calinski_harabasz_idx = calinski_harabasz_score(X, labels)

print(f'Cluster Evaluation for K = {n_clusters}:')
print(f'Silhouette Score: {silhouette_avg:.3f} (Higher is better, range -1 to 1)')
print(f'Davies-Bouldin Index: {davies_bouldin_idx:.3f} (Lower is better, minimum 0)')
print(f'Calinski-Harabasz Index: {calinski_harabasz_idx:.3f} (Higher is better)')

Cluster Evaluation for K = 6:
Silhouette Score: 0.168 (Higher is better, range -1 to 1)
Davies-Bouldin Index: 1.879 (Lower is better, minimum 0)
Calinski-Harabasz Index: 550.504 (Higher is better)
```

```
In [515]: pca = PCA(n_components=2, random_state=1)
pca_components = pca.fit_transform(X)

plt.figure(figsize=(10, 7))
plt.scatter(pca_components[:, 0], pca_components[:, 1], c=labels, cmap='viridis')
plt.title(f'Agglomerative Clusters (PCA, k={n_clusters})')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster Label')
plt.grid(True)
```



The k=6 plot reveals a slightly finer granularity on the left side, where the yellow and green clusters appear to differentiate further, indicating the additional cluster has likely refined the groupings within this dense region.

This re-segmentation on the left contrasts with the largely unchanged distribution of the dominant right-side cluster across both visualizations.